1. What are the fundamental building blocks of modern deep learning systems?
    a. Decision trees
    b. Support vector machines
    c. Neural networks
    d. Genetic algorithms

2. What do neural networks enable computers to do?
    a. Learn patterns
    b. Make decisions
    c. Understand human-like tasks
    d. All of the above

3. What is the arrangement of artificial neurons in neural networks?
    a. Clusters
    b. Layers
    c. Branches
    d. Circles

4. What is the process by which neural networks adjust internal weights and biases during training?
    1. Forward propagation
    2. Backpropagation
    3. Gradient descent
    4. Linear regression

5. How are neural networks categorized based on the complexity and number of hidden layers?
    a. Simple and complex
    b. Fast and slow
    c. Shallow and deep
    d. Linear and non-linear

6. How many hidden layers do shallow neural networks typically have?
    a. One or two
    b. Three to five
    c. More than five
    d. None

7. What type of problems are shallow neural networks suitable for?
    a. Complex image recognition
    b. Basic classification or regression
    c. Natural language modeling
    d. Speech understanding

8. What is a characteristic of deep neural networks (DNNs)?
    a. They have fewer hidden layers.
    b. They are ideal for simpler tasks.
    c. They learn hierarchical feature representations.
    d. They require less data.

9. Which of the following is NOT a type of neural network architecture?
    a. Feedforward Neural Networks (FNNs)
    b. Convolutional Neural Networks (CNNs)
    c. Recurrent Neural Networks (RNNs)
    d. Decision Trees

10. What is the basic architecture where data flows in one direction?
    a. Convolutional Neural Networks (CNNs)
    b. Recurrent Neural Networks (RNNs)
    c. Feedforward Neural Networks (FNNs)
    d. Long Short-Term Memory (LSTM) Networks

11. Which neural network architecture is designed for image processing tasks using feature extraction layers?
    a. Feedforward Neural Networks (FNNs)
    b. Convolutional Neural Networks (CNNs)
    c. Recurrent Neural Networks (RNNs)
    d. Long Short-Term Memory (LSTM) Networks

12. Which neural network architecture is suitable for sequential data, such as time series and text?
    a. Feedforward Neural Networks (FNNs)
    b. Convolutional Neural Networks (CNNs)
    c. Recurrent Neural Networks (RNNs)
    d. Long Short-Term Memory (LSTM) Networks

13. Which neural network architecture is an advanced RNN variant that handles long-term dependencies?
    a. Feedforward Neural Networks (FNNs)
    b. Convolutional Neural Networks (CNNs)
    c. Recurrent Neural Networks (RNNs)
    d. Long Short-Term Memory (LSTM) Networks

14. Which neural network architecture is used in modern natural language processing models like BERT and GPT?
    a. Feedforward Neural Networks (FNNs)
    b. Convolutional Neural Networks (CNNs)
    c. Recurrent Neural Networks (RNNs)
    d. Transformers

15. Who introduced the Perceptron?
    a. Geoffrey Hinton
    b. Yann LeCun
    c. Frank Rosenblatt
    d. Yoshua Bengio

16. What is the primary function of a Perceptron?
    a. To process images
    b. To map input features to an output
    c. To understand human language
    d. To generate new data

17. What are the learnable parameters in a Perceptron?
    a. Inputs and outputs
    b. Weights and bias
    c. Layers and nodes
    d. Activation functions

18. What type of output does a Perceptron produce?
    a. Continuous
    b. Binary
    c. Categorical
    d. Probabilistic

19. What are the components of a Perceptron?
    a. Input values, weights, bias, summation function, activation function
    b. Layers, nodes, connections
    c. Features, labels, predictions
    d. Training data, validation data, test data

20. What does the bias in a Perceptron do?
    a. Multiplies the input values
    b. Determines the importance of input features
    c. Shifts the activation function
    d. Calculates the sum of weighted inputs

21. What is the role of the activation function in a Perceptron?
    a. To calculate the weights
    b. To introduce non-linearity
    c. To sum the inputs
    d. To determine the bias

22. What is a limitation of a single-layer perceptron?
    a. It can only solve non-linear problems.
    b. It requires large amounts of data.
    c. It can only solve linearly separable problems.
    d. It has too many hidden layers.

23. What is a key enhancement of Multi-Layer Perceptrons (MLPs) compared to single-layer perceptrons?
    a. They have fewer layers.
    b. They can model non-linear decision boundaries.
    c. They use linear activation functions.
    d. They are simpler to train.

24. What algorithm is used to train Multi-Layer Perceptrons (MLPs)?
    a. Linear regression
    b. Backpropagation
    c. Support vector machine
    d. K-means clustering

25. What theorem states that an MLP with at least one hidden layer with enough neurons can approximate any continuous function?
    a. Central Limit Theorem
    b. Universal Approximation Theorem
    c. No Free Lunch Theorem
    d. Bayes' Theorem

26. What does a loss function measure?
    a. Model accuracy
    b. Model complexity
    c. The difference between predicted and actual values
    d. Training speed

27. What is the primary purpose of a loss function?
    a. To optimize hardware performance
    b. To guide the training process by quantifying errors
    c. To define the neural network architecture
    d. To preprocess input data

28. Which type of loss function is commonly used for regression problems?
    a. Cross-entropy loss
    b. Mean squared error (MSE)
    c. Binary cross-entropy
    d. Categorical cross-entropy

29. Which type of loss function is commonly used for binary classification problems?
    a. Mean squared error (MSE)
    b. Categorical cross-entropy
    c. Binary cross-entropy
    d. Hinge loss

30. Which type of loss function is commonly used for multi-class classification problems?
    a. Mean squared error (MSE)
    b. Binary cross-entropy
    c. Categorical cross-entropy
    d. Hinge loss

31. What is the role of optimization techniques in deep learning?
    a. To define the model architecture
    b. To minimize the loss function
    c. To preprocess data
    d. To deploy the model

32. Why do optimizers matter in training neural networks?
    a. They increase the training speed
    b. They ensure the model generalizes well
    c. They find the best set of weights to minimize the loss
    d. They prevent overfitting

33. Which of the following is a common type of optimizer?
    a. ReLU
    b. Sigmoid
    c. Adam
    d. Softmax

34. What is a characteristic of the Adam optimizer?
    a. It has a constant learning rate.
    b. It adapts the learning rate for each parameter.
    c. It only uses first-order moments.
    d. It is very slow.

35. What is overfitting in deep learning?
    a. When a model performs poorly on training data
    b. When a model performs well on training data but poorly on unseen data
    c. When a model generalizes perfectly to new data
    d. When a model trains for too short a time

36. How can overfitting be recognized?
    a. Large gap between training and validation performance
    b. High training accuracy and low validation accuracy
    c. Both training and validation accuracy are low
    d. Both training and validation accuracy are high

37. Which of the following is a strategy to prevent overfitting?
    a. Increasing the amount of training data
    b. Using more complex models
    c. Reducing the number of training epochs
    d. Early stopping

38. What is underfitting in deep learning?
    a. When a model performs well on training data
    b. When a model performs poorly on both training and unseen data
    c. When a model generalizes perfectly
    d. When a model trains for too long

39. How can underfitting be recognized?
    a. Large gap between training and validation performance
    b. High training accuracy and low validation accuracy
    c. Both training and validation accuracy are low
    d. High training accuracy and high validation accuracy

40. Which of the following is a strategy to fix underfitting?
    a. Reducing model complexity
    b. Increasing the amount of training data
    c. Using simpler models
    d. Early stopping

41. What does the bias-variance tradeoff describe?
    a. The relationship between training speed and accuracy
    b. The balance between model complexity and generalization
    c. The choice of activation function
    d. The selection of optimizer

42. What is hyperparameter tuning in deep learning?
    a. Adjusting the model's weights during training
    b. Selecting the best loss function
    c. Choosing the optimal values for parameters that control the learning process
    d. Preprocessing the input data

43. Which of the following is a key hyperparameter in a neural network?
    a. Weights
    b. Bias
    c. Learning rate
    d. Activations

44. Which of the following is a hyperparameter tuning method?
    a. Gradient descent
    b. Backpropagation
    c. Grid search
    d. Forward propagation

45. What is the first step in building a deep neural network?
    a. Training the model
    b. Designing the network architecture
    c. Compiling the model
    d. Evaluating performance

46. What does designing the network architecture involve?
    a. Choosing the loss function
    b. Selecting the optimizer
    c. Deciding on the number of layers, types of layers, and their connections
    d. Preprocessing the data

47. In the example architecture for MNIST digit classification, what type of layers are typically used?
    a. Recurrent layers
    b. Convolutional and dense layers
    c. Only dense layers
    d. Only recurrent layers

48. What does training a deep neural network involve?
    a. Defining the model architecture
    b. Selecting hyperparameters
    c. Feeding data through the network and updating weights using backpropagation
    d. Evaluating the model's performance

49. What does monitoring performance during training help with?
    a. Choosing the activation function
    b. Detecting and correcting training issues
    c. Defining the loss function
    d. Selecting the optimizer

50. What is early stopping?
    a. A technique to speed up training
    b. A method to prevent overfitting by halting training when validation performance stops improving
    c. A way to visualize model performance
    d. A type of optimizer